

AI OTHERWISE: RETHINKING WHAT AI SYSTEMS ARE ASSUMED TO BE

ICLR 2027 Workshop Proposal

Siya Qi
King's College London
[email]

Yoonseo Choi
KAIST
[email]

[Organizer 3]
[Institution]
[email]

[Organizer 4]
[Institution]
[email]

[Organizer 5]
[Institution]
[email]

ABSTRACT

Most current foundation models are built in the same way: next-token pretraining at scale, post-training on human preference, long reasoning before answering, and long context with persistent memory. Each of these choices improves average-case benchmarks, so each has become a default. These defaults work well for typical tasks, but they also limit the kinds of models that researchers consider. *AI Otherwise* studies what happens when a default is removed instead of improved. We call a model or system in which one default has been deliberately removed, inverted, or replaced an *alternative setting*. The workshop asks for which tasks such a setting is useful, how to build it without affecting the rest of the model, and how to evaluate its effects. The call has three themes: tasks that need a different model, methods for building alternative settings, and the impact of these systems on people, applications, and society. Only the third theme focuses on users. The program includes research papers, tiny papers, and working artifacts, as well as a hands-on session in which participants use alternative settings. The outcomes are an open catalogue of alternative settings and a synthesis paper written with participants.

1 WORKSHOP SUMMARY

The defaults of contemporary AI. Recent progress in language models (LM) has been accompanied by a growing set of shared assumptions about what AI systems should be like. Scaling laws showed that loss decreases predictably with model size, data, and compute (Kaplan et al., 2020; Hoffmann et al., 2022). Large pretrained models also turned out to have broad capabilities, and post-training on demonstrations and preference judgments turned them into general-purpose assistants (Ouyang et al., 2022; Bai et al., 2022). As these systems have become more capable, several design choices have increasingly come to be taken for granted: models should generate extended reasoning before answering (Wei et al., 2022; Guo et al., 2025), process longer contexts (Gemini Team, 2024), retain information across sessions (Packer et al., 2023), and act as agents that call tools until a task is completed (Yao et al., 2023). Most frontier models now embody some version of these choices. We call them the **defaults** of contemporary AI: choices that a new model is expected to adopt unless its developers deliberately decide otherwise. They are typically justified by the same kinds of evidence: improvements on aggregate benchmarks or increases in average human preference. As a result, they are often treated not as design choices, but as natural properties of what a capable AI system should be.

What the defaults cost. A default that improves average performance can still be a poor choice in specific uses. For example, training on human preference may make models agree with users even when the user is wrong (Sharma et al., 2024). It also reduces the diversity of model outputs (Kirk

Table 1: Five example alternative settings. Text-free AI is illustrated by an existing system, Jev. The organizers will build and release the other four. Each setting changes one default at a different level of the system. The last column lists hypotheses, not results.

Setting	Default removed (where)	How it can be built	Tasks it may suit, and the open question
<i>Text-free AI</i> returns a typed decision instead of generated text; e.g., Jev	Free-text generation (output format)	Outputs constrained to a supplied schema of options or score levels, each returned with a probability estimate	High-volume choice, scoring, and yes-or-no decisions consumed by software. Which tasks can be reduced to a typed decision without losing what the text carried?
<i>Goldfish AI</i> keeps only the last few turns	Long context and persistent memory (memory)	Truncated context, memory disabled, or targeted unlearning	Privacy-sensitive and one-off consultations. Which other capabilities degrade once memory is gone?
<i>Questioning AI</i> challenges the user’s premise, argument, or decision	Accommodating the user (post-training objective)	Preference data and rewards that favor well-founded challenges; steering against the sycophancy direction (Rimsky et al., 2024)	Reviewing research ideas, plans, and decisions; detecting false premises. How can a model challenge only when a challenge is warranted, and how is that evaluated?
<i>Read-only AI</i> observes but cannot act	Permission to act (agent harness)	A tool set restricted to reading, with no write or execute rights	Auditing and monitoring in high-stakes environments. How much of an agent benchmark can a read-only agent still complete?
<i>Long-tail AI</i> answers from the tail of its output distribution	Returning the most probable, most preferred answer (decoding)	Tail sampling, inverted preferences, or rewards against the mode	Brainstorming, hypothesis generation, and creative writing, where post-training has made outputs uniform. How can a rare but valuable answer be told apart from a wrong one?

et al., 2024) and of the text that users write with the model (Padmakumar & He, 2024). Long reasoning spends substantial computation on simple problems (Chen et al., 2025). Persistent memory raises privacy concerns for users (Chen et al., 2026). Most existing work addresses these problems while keeping the default setting, for example by reducing sycophancy, shortening reasoning, or adding memory controls. These works are useful, but they largely focus on refining the default rather than asking what happens when the default changes. This question has received less attention as a research focus.

Alternative settings. This workshop focuses on that question. We define an **alternative setting** as a model or system in which one default has been deliberately changed: removed, inverted, or replaced. The change can remove a capability, such as memory or the ability to act. It can also change a design choice in training, architecture, decoding, or the agent harness: the program around the model that controls its inputs and the tools it can call. The point is not simply to remove a capability, but to explore what changes when a familiar default is changed. Jev is one example.¹ It departs from the common training procedure of language models and does not generate free text, instead returning structured outputs with probability estimates for downstream software. Giving up text generation would be a loss for a general assistant, but an advantage for faster and more direct decision making in this setting. Table 1 lists potential research examples, each changing a default at a different level of the system. These examples are starting points for exploration, not a taxonomy.

¹Jev, [https://en.wikipedia.org/wiki/Jev_\(AI_model\)](https://en.wikipedia.org/wiki/Jev_(AI_model)).

Problems we are dealing with. The default costs we discussed point to a broader gap in how AI systems are explored. Research often starts from a relatively narrow set of model and system configurations and asks how to make them more capable, align, or useful. This leaves less room to systematically explore what other configurations might offer. A different setting may be better suited to a particular task, support a different form of human interaction, or reveal something about the capabilities and limitations of current models. We therefore ask a broader question: *what becomes possible when familiar defaults are changed?* The workshop aims to make progress on two questions. First, how can alternative settings in architecture, training objectives, decoding, memory, and agent harnesses be explored systematically? This requires changing one default and reporting results that other groups can reproduce and compare. Second, which alternative settings are useful for particular tasks or human goals, and what new insights can they reveal? One workshop may not answer all these questions. We expect it to produce initial answers by bringing together researchers who build models and researchers who study how models are used.

Why now, and why ICLR. This workshop is particularly relevant now for two reasons. First, frontier models increasingly follow similar pipeline, which leaves little room for alternatives in mainstream research. Second, the tools needed to build alternatives are now available. Refusal is mediated by a single direction in activation space (Arditi et al., 2024), and inference-time steering can change specific behaviors (Li et al., 2023; Rinsky et al., 2024). Fine-tuning is less expensive, and agent harnesses are easy to modify. Removing a single default is therefore feasible, but not yet reliable, because fine-tuning on a narrow behavior can change unrelated behaviors (Betley et al., 2026). The main technical problems of the workshop are changing one default without affecting the others, and evaluating models whose value standard benchmarks do not capture. These are machine learning problems, and ICLR brings together the relevant communities in representation learning, post-training, and agents. The workshop therefore starts from the model and treats human use as one part of the evaluation.

Why now, and why ICLR. This workshop is particularly relevant now for two reasons. First, frontier models increasingly share similar design choices, making it timely to explore what other model and system configurations can offer. Second, recent advances make such exploration increasingly feasible. Steering and fine-tuning provide practical ways to modify model behavior, while agent harnesses and other system components can be changed independently of the underlying model. These tools make it possible to construct alternative settings and study them systematically, while also raising an important methodological challenge: changing one default without unintentionally changing others. The workshop is therefore centered on questions about model behavior, training, inference, and system design, while also asking how these alternatives should be evaluated in use. This combination connects naturally to ICLR’s communities in representation learning, post-training, agents, and evaluation.

2 TOPICS OF INTEREST

We invite submissions on alternative settings. The three themes below describe the work we hope to receive. They are not exclusive, and settings not listed in Table 1 are welcome. Submissions do not need to be complete: we welcome ideas for new architectures or new ways of using models, preliminary experiments, short idea papers, demo designs, and user studies. Small observations from everyday use are as welcome as large studies.

Tasks that call for a different model. Papers that identify a specific task for which a default model performs poorly or is unsuitable, explain which default causes the problem, and describe the setting the task requires. A single well-defined task is sufficient. Papers that describe such a task without a system or experiment can be submitted as tiny papers. One example is decisions made inside software, which require a typed value rather than text.

Building alternative settings. Papers that remove one mainstream setting in pretraining, architecture, post-training, decoding, memory, or the agent harness. We also welcome comparisons of removing the same setting at different levels, steering and interpretability methods for removing or adding a default, and training data and objectives for specific behaviors.

Impact of alternative AI. Papers discussing how alternative settings affect user behaviors, applications, and broader society. We welcome user studies, comparisons of the same task under a default model and an alternative setting, multi-session evaluation protocols, negative results, and discussion of how to prevent harmful effects.

3 RELATION TO PREVIOUS WORKSHOPS

This is the first edition of **AI Otherwise**. Several recent workshops have examined properties that we treat as defaults, but from different perspectives. Some focus on improving these properties. Workshop on memory for LLM agents (**Memory for LLM-Based Agentic Systems**, ICLR 2026) studies how memory can be made more capable or safer, and work on pluralistic alignment (**Pluralistic Alignment**, ICML 2026) studies how models can better represent preferences. Other workshops examine the assumptions and limitations of existing systems, including failure modes of current models (**I Can’t Believe It’s Not Better**, ICLR 2025 and 2026), the assumptions behind learning from human feedback (**Models of Human Feedback for AI Alignment**, ICML 2024 and 2025), and when people should rely on language models (**ORIGen**, COLM 2025). Work on mutual adaptation between models and people (**Bidirectional Human-AI Alignment**, ICLR 2025) is particularly close to our interests, but focuses on improving alignment between them.

Our focus is complementary to these efforts. Rather than assuming that a capability or design choice should be improved, we ask what can be learned by deliberately changing it and studying the resulting system. In particular, **I Can’t Believe It’s Not Better** shares our interest in questioning whether the dominant direction of AI development is necessarily better. We take this question in a constructive direction: rather than only identifying what does not work in current systems, we ask what other configurations we can build and what they enable or reveal. Existing work studies many individual properties in isolation, but there is less work that brings them together as alternative configurations of AI systems and studies their task, human, and scientific value within a common framework.

AI Otherwise therefore complements rather than replaces these lines of work: it provides a shared setting in which researchers can build, compare, and study AI systems in which familiar defaults have been deliberately changed.

4 WORKSHOP FORMAT AND TENTATIVE SCHEDULE

The workshop is a one-day, in-person event (Table 2) with four invited talks, six contributed talks, two poster sessions, and a panel. It also includes an *encounter session*. Participants move between stations and try the artifacts submitted by authors and the four systems in Table 1, spending a few minutes on each. They record on cards what they expected, what they did differently, and what the system was useful for. The cards are placed on a *settings map*, a board with one column per default that stays up for the whole day. In the afternoon, small groups use the map to summarize, for each setting, the task it served and the open questions. The panel takes its questions from these summaries. The encounter session gives all participants a shared experience to discuss, and the map records the results for the synthesis paper. The morning covers tasks and methods for building alternative settings, and the afternoon covers evaluation and impact. Organizers only give the opening and closing remarks and moderate the panel.

5 INVITED SPEAKERS AND PANELISTS

We plan four invited talks of 30 minutes each, following the three themes of the call. Three speakers work on machine learning, covering post-training or steering, tasks and models outside the standard pipeline, and evaluation. The fourth works on human–AI interaction and covers the impact theme. All invited speakers also serve as panelists. *[To be filled: names, affiliations, OpenReview IDs, and confirmation status for each speaker. Each speaker may be listed on at most two ICLR 2027 workshop proposals.]*

Table 2: Tentative schedule (single day, in-person). Times will be adjusted to the ICLR 2027 workshop-day template.

Time	Activity
<i>Morning: Tasks and alternative settings</i>	
08:50–09:00	Opening remarks: the defaults, the four systems, and the settings map.
09:00–09:30	Invited talk 1 (removing a default: post-training, steering, and the harness).
09:30–10:00	Invited talk 2 (tasks the defaults do not serve, and models built for them).
10:00–10:40	Contributed talks I: three 12-minute talks.
10:40–11:20	Coffee break and poster session I.
11:20–12:15	Encounter session. Participants rotate through stations run by artifact authors and by the organizers with the four systems, and record their observations on cards for the settings map.
12:15–13:30	Lunch.
<i>Afternoon: Evaluation and impact</i>	
13:30–14:00	Invited talk 3 (evaluating a model whose value standard benchmarks do not measure).
14:00–14:30	Invited talk 4 (how people and applications change when a default is removed).
14:30–15:10	Contributed talks II: three 12-minute talks.
15:10–15:50	Coffee break and poster session II.
15:50–16:30	Reflection in small groups. Groups of 5–6 cluster the cards on the settings map and write up, for each setting, the task it served, what changed in how it was used, and what remains open.
16:30–17:15	Panel: <i>Which defaults are worth keeping, and for whom?</i> Invited speakers and one contributed-paper author, moderated by an organizer.
17:15–17:30	Closing remarks, best paper and best artifact awards, next steps.

6 CONTRIBUTED WORK

All accepted contributions receive a poster slot, six are selected for talks, and accepted artifacts are presented hands-on in the encounter session. The workshop is non-archival, and accepted work is publicly visible on OpenReview. Titles of accepted talks will be posted on the website before the conference. The call will state that we prioritize new and ongoing work over work already published at a machine learning venue. We offer three tracks. **Research papers** of up to 5 pages in the ICLR format, excluding references, present empirical, theoretical, or methodological work on any of the themes; work under review elsewhere is welcome. **Tiny papers** of up to 2 pages, following the ICLR guidance for late-breaking and less formal work, may report a single experiment, a negative result, a well-argued question, or a single task for which a default was the wrong choice. **Artifacts** are working or partially working alternative settings that participants can use, submitted as a 2-page description of the removed default, the method, and observations, with a link to the running system or a video. Authors bring their own devices to the encounter session. Early-stage, playful, and unconventional systems are encouraged.

LLM usage policy. We follow the ICLR 2027 policy on large language model usage. Language models may assist authors and reviewers, but they may not serve as primary authors or reviewers, and AI-generated submissions are not permitted in the tiny-paper track.

Review process and timeline. Papers and tiny papers are reviewed double-blind by at least two reviewers from a program committee of 30–40 researchers drawn from post-training, steering and interpretability, evaluation, agents, and human–AI interaction. Artifacts are reviewed single-blind, since a running system often identifies its authors. Conflicts of interest are managed through OpenReview, and organizers do not handle submissions from their own institutions. We anticipate 40–60 submissions and an acceptance rate of roughly 60–70%, given that the workshop is non-archival. Best paper and best artifact awards are selected by the program committee. Tentative timeline: call published by mid-December 2026, submission deadline January 29, 2027, reviews due February 19, author notification, OpenReview release, and import to iclr.cc by February 26, camera-ready March 19, and schedule finalized by March 26 on iclr.cc and April 14 in the organizer portal.

7 ORGANIZERS

Most organizers work in machine learning, on LLM reliability and interpretability, post-training and alignment, and evaluation. The team also includes expertise in human–AI interaction for the impact theme. *[3–5 sentences per organizer: institution, research focus, relevant publications, OpenReview ID, and other ICLR 2027 proposals they are involved in. Up to 8 organizers.]* **Siya Qi** is a Ph.D. candidate in natural language processing at King’s College London. Her research focuses on the reliability of large language models and their real-world applications, including hallucination, LLM agents, theory of mind, and mechanistic interpretability. *[Relevant publications, OpenReview ID, other ICLR 2027 proposals.]* Lead organizer, responsible for the program, the call, and the synthesis. **Yoonseo Choi** is a Ph.D. candidate at KAIST, working with Prof. Juho Kim. She studies how user perspectives can be represented, interpreted, and integrated throughout the lifecycle of AI-mediated systems, with interests in human–AI interaction, generative proxies, creator ecosystems, and responsible AI. *[Relevant publications, OpenReview ID, other ICLR 2027 proposals.]* Leads the encounter and reflection sessions. **[Organizer 3]** *[Institution, bio: post-training, alignment, or steering.]* Responsible for the four systems and the artifact track. **[Organizer 4]** *[Institution, bio: evaluation or studies of deployed models in use.]* Leads the program committee and review process. **[Organizer 5]** *[Institution, bio: senior researcher.]* Advisory role, responsible for speaker invitations and the panel.

Organizational experience. *[To be filled.]*

Nominated reviewer. Following the ICLR 2027 requirement, we nominate **[Name]**, **[Affiliation]**, a mid-career/senior researcher who is not an organizer, to serve on the workshop proposal review committee. *[OpenReview ID.]*

8 ANTICIPATED AUDIENCE, OUTREACH, AND DIVERSITY

We anticipate 100–150 attendees and 40–60 submissions, based on the attendance of recent ICLR workshops on alignment, agents, and evaluation. The encounter session scales by adding stations. We will publish a website with the call, the four systems and their recipes, and the accepted work, and distribute the call through ICLR and OpenReview channels, machine learning mailing lists, social media, and direct invitations to the groups whose work is cited in Section 1. The tiny-paper and artifact tracks lower the barrier for early-stage work and for researchers who have built an alternative setting but have not yet submitted to ICLR.

We commit to diversity of viewpoint, since the workshop brings together subfields that rarely share a venue, and to diversity of gender, ethnicity, seniority, affiliation type, and geography across organizers, speakers, panelists, and the program committee. The organizing team *[spans N institutions in N countries, includes N women, and includes senior and early-career researchers]*. We reserve one contributed-talk slot and one panel seat for an early-career author and will seek ICLR travel support for student presenters.

9 VIRTUAL ACCESS AND OUTCOMES

ICLR 2027 workshops are in-person. With speaker consent, we will record the invited talks and the panel and publish them on the workshop website, publish all accepted papers and artifact descriptions on OpenReview, and display posters and pre-recorded videos for authors who cannot attend due to exceptional circumstances. The four systems will be released with their prompts, steering vectors, or fine-tuning recipes so that they can be run and extended remotely.

The workshop has two outcomes, each linked to a question in Section 1. The first is an open catalogue of alternative settings. Each entry records the default that is removed, the level at which it is removed, the target task, the recipe, and observations from use. Because each entry includes a recipe and a task, results can be reproduced and compared, which addresses the first question. The tasks collected under the first theme address the second question. The second outcome is a synthesis paper written with participants. It will use the settings map and the panel discussion to propose shared terminology and a research agenda.

Safety and ethics. Removing a default must not remove a model’s safety behavior. Submissions that aim to disable safeguards are out of scope. Two settings need extra care. A model trained to challenge the user can be uncomfortable to use, and a model that samples from the tail of its distribution is more likely to produce wrong or misleading content. The organizers will review all artifacts before the encounter session. Each artifact is introduced with a short description of how it differs from a standard model, and participants can stop at any time. Artifacts that log participant data must state this in the submission, obtain consent at the station, and store no identifiable data. Papers that report user studies should state their ethics approval where applicable. The workshop itself collects no personal data except anonymized photographs of the settings map. The four systems released by the organizers will be documented with their intended use and known risks.

REFERENCES

- Andy Arditi, Oscar Obeso, Aaquib Syed, Daniel Paleka, Nina Panickssery, Wes Gurnee, and Neel Nanda. Refusal in language models is mediated by a single direction. In *Advances in Neural Information Processing Systems 37 (NeurIPS 2024)*, 2024.
- Yuntao Bai, Andy Jones, Kamal Ndousse, et al. Training a helpful and harmless assistant with reinforcement learning from human feedback, 2022. URL <https://arxiv.org/abs/2204.05862>.
- Jan Betley, Niels Warncke, Anna Sztyber-Betley, Daniel Tan, Xuchan Bao, Martín Soto, Megha Srivastava, Nathan Labenz, and Owain Evans. Training large language models on narrow tasks can lead to broad misalignment. *Nature*, 649(8097):584–589, 2026. doi: 10.1038/s41586-025-09937-5.
- Cheng Chen, Maria D. Molina, Mengqi Liao, and Eugene Cho Snyder. Relational gains, privacy strains: Exploring users’ perceptions and experiences with ChatGPT’s memory feature. In *Proceedings of the 2026 CHI Conference on Human Factors in Computing Systems (CHI ’26)*. ACM, 2026. doi: 10.1145/3772318.3791635.
- Xingyu Chen, Jiahao Xu, Tian Liang, Zhiwei He, Jianhui Pang, Dian Yu, Linfeng Song, Qiuzhi Liu, Mengfei Zhou, Zhuosheng Zhang, Rui Wang, Zhaopeng Tu, Haitao Mi, and Dong Yu. Do NOT think that much for $2+3=?$ on the overthinking of long reasoning models. In *Proceedings of the 42nd International Conference on Machine Learning (ICML 2025)*, volume 267 of *Proceedings of Machine Learning Research*, pp. 9487–9499, 2025.
- Gemini Team. Gemini 1.5: Unlocking multimodal understanding across millions of tokens of context. *arXiv preprint arXiv:2403.05530*, 2024.
- Daya Guo, Dejian Yang, Haowei Zhang, et al. DeepSeek-R1 incentivizes reasoning in LLMs through reinforcement learning. *Nature*, 645:633–638, 2025. doi: 10.1038/s41586-025-09422-z.
- Jordan Hoffmann, Sebastian Borgeaud, Arthur Mensch, Elena Buchatskaya, Trevor Cai, Eliza Rutherford, Diego de Las Casas, Lisa Anne Hendricks, Johannes Welbl, Aidan Clark, et al. An empirical analysis of compute-optimal large language model training. In *Advances in Neural Information Processing Systems 35 (NeurIPS 2022)*, 2022.
- Jared Kaplan, Sam McCandlish, Tom Henighan, Tom B. Brown, Benjamin Chess, Rewon Child, Scott Gray, Alec Radford, Jeffrey Wu, and Dario Amodei. Scaling laws for neural language models. *arXiv preprint arXiv:2001.08361*, 2020.
- Robert Kirk, Ishita Mediratta, Christoforos Nalmpantis, Jelena Luketina, Eric Hambro, Edward Grefenstette, and Roberta Raileanu. Understanding the effects of RLHF on LLM generalisation and diversity. In *The Twelfth International Conference on Learning Representations (ICLR 2024)*, 2024. URL <https://openreview.net/forum?id=PXD3FAVHJT>.
- Kenneth Li, Oam Patel, Fernanda Viégas, Hanspeter Pfister, and Martin Wattenberg. Inference-time intervention: Eliciting truthful answers from a language model. In *Advances in Neural Information Processing Systems 36 (NeurIPS 2023)*, 2023.

- Long Ouyang, Jeff Wu, Xu Jiang, et al. Training language models to follow instructions with human feedback. In *Advances in Neural Information Processing Systems 35 (NeurIPS 2022)*, 2022.
- Charles Packer, Sarah Wooders, Kevin Lin, Vivian Fang, Shishir G. Patil, Ion Stoica, and Joseph E. Gonzalez. MemGPT: Towards LLMs as operating systems, 2023. URL <https://arxiv.org/abs/2310.08560>.
- Vishakh Padmakumar and He He. Does writing with language models reduce content diversity? In *The Twelfth International Conference on Learning Representations (ICLR 2024)*, 2024. URL <https://openreview.net/forum?id=Feiz5HtCD0>.
- Nina Rimskey, Nick Gabrieli, Julian Schulz, Meg Tong, Evan Hubinger, and Alexander Matt Turner. Steering Llama 2 via contrastive activation addition. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (ACL 2024)*, pp. 15504–15522, 2024. doi: 10.18653/v1/2024.acl-long.828.
- Mrinank Sharma, Meg Tong, Tomasz Korbak, et al. Towards understanding sycophancy in language models. In *The Twelfth International Conference on Learning Representations (ICLR 2024)*, 2024. URL <https://openreview.net/forum?id=tvhaxkMKAn>.
- Jason Wei, Xuezhi Wang, Dale Schuurmans, Maarten Bosma, Brian Ichter, Fei Xia, Ed H. Chi, Quoc V. Le, and Denny Zhou. Chain-of-thought prompting elicits reasoning in large language models. In *Advances in Neural Information Processing Systems 35 (NeurIPS 2022)*, 2022.
- Shunyu Yao, Jeffrey Zhao, Dian Yu, Nan Du, Izhak Shafran, Karthik Narasimhan, and Yuan Cao. ReAct: Synergizing reasoning and acting in language models. In *The Eleventh International Conference on Learning Representations (ICLR 2023)*, 2023. URL https://openreview.net/forum?id=WE_vluYUL-X.